

TP1 : Apache HBase

Rapport

Réalisé par :

Abderrahim ALAKOUCHE
Abdellah AGHLALOU

Encadré par :

Mme. Dounia ZAIDOUNI

PLAN

Introduction

I

Installation et configuration d'Apache Hbase

II

Manipulation de « HBase »

III

L'utilisation de l'API Java de HBase

IV

Traitement de données avec Spark



HBase est un système de gestion de base de données non-relationnelles distribué, écrit en Java, disposant d'un stockage structuré pour les grandes tables.[1]

Les objectifs de ce TP sont les suivants :

- Installation et configuration d'Apache HBase en mode « Standalone » et en mode « Pseudo distribué ».
- Manipulation de « HBase » (opérations create, put , get, list, scan, disable, drop, exists)
- L'utilisation de l'API Java de Hbase.
- Traitement de données avec Spark.

Etape1: Installation et configuration d'Apache Hbase :

```
agh@agh-VirtualBox: ~  
agh@agh-VirtualBox:~/Downloads$ ls  
hadoop-2.7.4.tar.gz  hbase-1.4.7-bin.tar.gz  spark-2.4.3-bin-hadoop2.7.tgz  
hbase-1.4.7         jdk-8u71-linux-x64.tar.gz  tp_hbase.pdf  
agh@agh-VirtualBox:~/Downloads$ mv hbase-1.4.7 hbase  
agh@agh-VirtualBox:~/Downloads$ ls  
hadoop-2.7.4.tar.gz  hbase-1.4.7-bin.tar.gz  spark-2.4.3-bin-hadoop2.7.tgz  
hbase                jdk-8u71-linux-x64.tar.gz  tp_hbase.pdf  
agh@agh-VirtualBox:~/Downloads$ sudo mv hbase /usr/local/  
[sudo] password for agh:  
agh@agh-VirtualBox:~/Downloads$ cd  
agh@agh-VirtualBox:~$ vim .bashrc  
agh@agh-VirtualBox:~$ vim .bashrc  
agh@agh-VirtualBox:~$ source .bashrc  
agh@agh-VirtualBox:~$
```

Le fichier de Hbase

Renommer le fichier

déplacement

Ajouter le /bin dans le
PATH

```
agh@agh-VirtualBox: ~  
# enable programmable completion features (you don't need to enable  
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile  
# sources /etc/bash.bashrc).  
if ! shopt -oq posix; then  
    if [ -f /usr/share/bash-completion/bash_completion ]; then  
        . /usr/share/bash-completion/bash_completion  
    elif [ -f /etc/bash_completion ]; then  
        . /etc/bash_completion  
    fi  
fi  
#hadoop variables start  
export JAVA_HOME=/opt/java/jdk1.8.0_71/  
export HADOOP_INSTALL=/usr/local/hadoop  
export PATH=$PATH:$HADOOP_INSTALL/bin  
export PATH=$PATH:$HADOOP_INSTALL/sbin  
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL  
export HADOOP_COMMON_HOME=$HADOOP_INSTALL  
export HADOOP_HDFS_HOME=$HADOOP_INSTALL  
export YARN_HOME=$HADOOP_INSTALL  
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native  
#export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"  
#HADOOP VARIABLES END  
export SPARK_HOME=/usr/local/spark  
export PATH=$PATH:$SPARK_HOME/bin  
export HBASE_HOME=/usr/local/hbase  
export PATH=$PATH:$HBASE_HOME/bin
```

```
agh@agh-VirtualBox: /usr/local/hbase/conf
agh@agh-VirtualBox:~$ java -version
java version "1.8.0_71"
Java(TM) SE Runtime Environment (build 1.8.0_71-b15)
Java HotSpot(TM) 64-Bit Server VM (build 25.71-b15, mixed mode)
agh@agh-VirtualBox:~$ cd /usr/local/hbase/conf/
agh@agh-VirtualBox:/usr/local/hbase/conf$ vim hbase env.sh
```

```
agh@agh-VirtualBox: /usr/local/hbase/conf
# * See the License for the specific language governing permissions and
# * limitations under the License.
# */

# Set environment variables here.

# This script sets variables multiple times over the course of starting an hbase
# process,
# so try to keep things idempotent unless you want to take an even deeper look
# into the startup scripts (bin/hbase, etc.)

# The java implementation to use.  Java 1.7+ required.
export JAVA_HOME=/opt/java/jdk1.8.0_71
# Extra Java CLASSPATH elements.  Optional.
# export HBASE_CLASSPATH=

# The maximum amount of heap to use. Default is left to JVM default.
# export HBASE_HEAPSIZE=1G

# Uncomment below if you intend to use off heap cache. For example, to allocate
# 8G of
# offheap, set the value to "8G".
# export HBASE_OFFHEAPSIZE=1G

# Extra Java runtime options.
# Below are what we set by default.  May only work with SUN JVM.
# For more on why as well as other possible settings,
# see http://wiki.apache.org/hadoop/PerformanceTuning
```

```
agh@agh-VirtualBox:/usr/local/hbase/conf$ cd
agh@agh-VirtualBox:~$ sudo vim /etc/hosts
[sudo] password for agh:
agh@agh-VirtualBox:~$ sudo reboot
```

```
127.0.0.1    localhost
127.0.0.1    agh-VirtualBox

# The following lines are desirable for IPv6 capable hosts
::1         ip6-localhost ip6-loopback
fe00::0     ip6-localnet
ff00::0     ip6-mcastprefix
ff02::1     ip6-allnodes
ff02::2     ip6-allrouters
```

Etape2:Démarrage du cluster Hadoop configuré dans notre machine .

```
agh@agh-VirtualBox:/usr/local/hadoop_store$ rm -rf *
agh@agh-VirtualBox:/usr/local/hadoop_store$ ls
agh@agh-VirtualBox:/usr/local/hadoop_store$ ls
agh@agh-VirtualBox:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hd
fs/namenode
agh@agh-VirtualBox:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hd
fs/datanode
agh@agh-VirtualBox:/usr/local/hadoop_store$ chown -R agh /usr/local/hadoop_sto
re/hdfs/namenode
agh@agh-VirtualBox:/usr/local/hadoop_store$ chown -R agh /usr/local/hadoop_sto
re/hdfs/datanode
agh@agh-VirtualBox:/usr/local/hadoop_store$
```


Formatage de namenode :

```
agh@agh-VirtualBox: /usr/local/hadoop/etc/hadoop
21/02/14 10:39:23 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.windows
.minutes = 1,5,25
21/02/14 10:39:23 INFO namenode.FSNamesystem: Retry cache on namenode is enable
d
21/02/14 10:39:23 INFO namenode.FSNamesystem: Retry cache will use 0.03 of tota
l heap and retry cache entry expiry time is 600000 millis
21/02/14 10:39:23 INFO util.GSet: Computing capacity for map NameNodeRetryCache
21/02/14 10:39:23 INFO util.GSet: VM type          = 64-bit
21/02/14 10:39:23 INFO util.GSet: 0.029999999329447746% max memory 966.7 MB = 2
97.0 KB
21/02/14 10:39:23 INFO util.GSet: capacity          = 2^15 = 32768 entries
21/02/14 10:39:23 INFO namenode.FSImage: Allocated new BlockPoolId: BP-80703196
5-127.0.0.1-1613299163364
21/02/14 10:39:23 INFO common.Storage: Storage directory /usr/local/hadoop_stor
e/hdfs/namenode has been successfully formatted.
21/02/14 10:39:23 INFO namenode.FSImageFormatProtobuf: Saving image file /usr/l
ocal/hadoop_store/hdfs/namenode/current/fsimage.ckpt_00000000000000000000 using
no compression
21/02/14 10:39:23 INFO namenode.FSImageFormatProtobuf: Image file /usr/local/ha
doop_store/hdfs/namenode/current/fsimage.ckpt_00000000000000000000 of size 320 b
ytes saved in 0 seconds.
21/02/14 10:39:24 INFO namenode.NNStorageRetentionManager: Going to retain 1 im
ages with txid >= 0
21/02/14 10:39:24 INFO util.ExitUtil: Exiting with status 0
21/02/14 10:39:24 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at localhost/127.0.0.1
*****/
agh@agh-VirtualBox: /usr/local/hadoop/etc/hadoop$
```

Lancement de namenode, datanode, secondary namenode , yarn

```
agh@agh-VirtualBox: /usr/local/hadoop/etc/hadoop
agh@agh-VirtualBox:/usr/local/hadoop/etc/hadoop$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
21/02/14 10:40:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-agh-namenode-agh-VirtualBox.out
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-agh-datanode-agh-VirtualBox.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-agh-secondarynamenode-agh-VirtualBox.out
21/02/14 10:41:10 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-agh-resourcemanager-agh-VirtualBox.out
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-agh-nodemanager-agh-VirtualBox.out
agh@agh-VirtualBox:/usr/local/hadoop/etc/hadoop$
```

Vérification

```
-----  
Live datanodes (1):  
  
Name: 127.0.0.1:50010 (localhost)  
Hostname: localhost  
Decommission Status : Normal  
Configured Capacity: 31147880448 (29.01 GB)  
DFS Used: 24576 (24 KB)  
Non DFS Used: 11531415552 (10.74 GB)  
DFS Remaining: 18010599424 (16.77 GB)  
DFS Used%: 0.00%  
DFS Remaining%: 57.82%  
Configured Cache Capacity: 0 (0 B)  
Cache Used: 0 (0 B)  
Cache Remaining: 0 (0 B)  
Cache Used%: 100.00%  
Cache Remaining%: 0.00%  
Xceivers: 1  
Last contact: Sun Feb 14 10:42:06 WET 2021  
  
agh@agh-VirtualBox:/usr/local/hadoop/etc/hadoop$
```

configuration d'Apache Hbase pour un mode « Standalone » :

```
agh@agh-VirtualBox:/usr/local/hadoop/etc/hadoop$ cd /usr/local/hbase/conf/
agh@agh-VirtualBox:/usr/local/hbase/conf$ vim hbase site.xml
agh@agh-VirtualBox:/usr/local/hbase/conf$ ls
hadoop-metrics2-hbase.properties  hbase-policy.xml  regionserver
hbase-env.cmd                    hbase-site.xml
hbase-env.sh                     log4j.properties
agh@agh-VirtualBox:/usr/local/hbase/conf$ vim hbase-site.xml
```

```
* regarding copyright ownership. The ASF licenses this file
* to you under the Apache License, Version 2.0 (the
* "License"); you may not use this file except in compliance
* with the License. You may obtain a copy of the License at
*
* http://www.apache.org/licenses/LICENSE-2.0
*
* Unless required by applicable law or agreed to in writing, software
* distributed under the License is distributed on an "AS IS" BASIS,
* WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
* See the License for the specific language governing permissions and
* limitations under the License.
*/
-->
<configuration>
  <property>
    <name>hbase.rootdir</name>
    <value>file:///home/agh/hbase</value>
  </property>
  <property>
    <name>hbase.zookeeper.property.dataDir</name>
    <value>/home/agh/zookeeper</value>
  </property>
  <property>
    <name>hbase.unsafe.stream.capability.enforce</name>
    <value>>false</value>
  </property>
</configuration>
```

Lancement de Hbase

```
agh@agh-VirtualBox:/usr/local/hbase$ jps
4469 Jps
4181 HMaster
2567 SecondaryNameNode
2379 DataNode
2254 NameNode
2719 ResourceManager
agh@agh-VirtualBox:/usr/local/hbase$
```

```
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.7, r763f27f583cf8fd7ecf79fb6f3ef57f1615dbf9b, Tue Aug 28 14:40:11 P
DT 2018

hbase(main):001:0> status
1 active master, 0 backup masters, 1 servers, 0 dead, 2.0000 average load

hbase(main):002:0>
```


- Création des fichiers nécessaires pour Hbase dans le HDFS :

```
agh@agh-VirtualBox:/usr/local/hbase$ hadoop fs -mkdir -p /hbase
21/02/14 11:03:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
agh@agh-VirtualBox:/usr/local/hbase$ hadoop fs -mkdir -p /agh/zookeeper
21/02/14 11:04:43 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
agh@agh-VirtualBox:/usr/local/hbase$ hadoop fs -ls /
21/02/14 11:05:20 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - agh supergroup          0 2021-02-14 11:04 /agh
drwxr-xr-x  - agh supergroup          0 2021-02-14 11:03 /hbase
agh@agh-VirtualBox:/usr/local/hbase$ hadoop fs -ls /agh
21/02/14 11:05:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x  - agh supergroup          0 2021-02-14 11:04 /agh/zookeeper
agh@agh-VirtualBox:/usr/local/hbase$
```

```
* regarding copyright ownership. The ASF licenses this file
* to you under the Apache License, Version 2.0 (the
* "License"); you may not use this file except in compliance
* with the License. You may obtain a copy of the License at
*
* http://www.apache.org/licenses/LICENSE-2.0
*
* Unless required by applicable law or agreed to in writing, software
* distributed under the License is distributed on an "AS IS" BASIS,
* WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
* See the License for the specific language governing permissions and
* limitations under the License.
*/
-->
<configuration>
  <property>
    <name>hbase.rootdir</name>
    <value>hdfs://localhost:54310/hbase</value>
  </property>
  <property>
    <name>hbase.zookeeper.property.dataDir</name>
    <value>hdfs://localhost:54310/agh/zookeeper</value>
  </property>
  <property>
    <name>hbase.cluster.distributed</name>
    <value>true</value>
  </property>
</configuration>
```

Lancement de Hbase :

```
agh@agh-VirtualBox:/usr/local/hbase$ ./bin/start-hbase.sh
localhost: running zookeeper, logging to /usr/local/hbase/bin/../../logs/hbase-agh-zookeeper-agh-VirtualBox.out
running master, logging to /usr/local/hbase/logs/hbase-agh-master-agh-VirtualBox.out
: running regionserver, logging to /usr/local/hbase/logs/hbase-agh-regionserver-agh-VirtualBox.out
agh@agh-VirtualBox:/usr/local/hbase$ jps
5330 Jps
2567 SecondaryNameNode
5096 HQuorumPeer ✓
2379 DataNode
2254 NameNode
5150 HMaster ✓
2719 ResourceManager
5215 HRegionServer ✓
agh@agh-VirtualBox:/usr/local/hbase$
```


-Requête de création d'une base de donne avec Hbase shell
-la commande list pour lister les tables existants

```
hbase(main):001:0> create 'registre_ventes','client','ventes'
0 row(s) in 4.2580 seconds

=> Hbase::Table - registre_ventes
hbase(main):002:0> list
TABLE
registre_ventes
1 row(s) in 0.0920 seconds

=> ["registre_ventes"]
hbase(main):003:0> █
```

Ajout d'un nouveau cell :

Put 'nom_table', ' num_of_cell', 'name_columnfamily:name_column', 'valeur'

```
agh@agh-VirtualBox: /usr/local/hbase/bin
hbase(main):019:0> scan 'registre_ventes'
ROW          COLUMN+CELL
 101         column=client:nom, timestamp=1613303968964, value=Mohamed
              A
 101         column=client:ville, timestamp=1613303969162, value=Casabl
              anca
 101         column=ventes:montant, timestamp=1613303969530, value=2000
              ,00 MAD
 101         column=ventes:produit, timestamp=1613303969357, value=Chai
              ses
 102         column=client:nom, timestamp=1613303969673, value=Amine B
 102         column=client:ville, timestamp=1613303969869, value=Sal\xC
              3\xA9
 102         column=ventes:montant, timestamp=1613303970142, value=300,
              00 MAD
 102         column=ventes:produit, timestamp=1613303970001, value=Lamp
              es
 103         column=client:nom, timestamp=1613303970289, value=Younes C
 103         column=client:ville, timestamp=1613303970380, value=Rabat
 103         column=ventes:montant, timestamp=1613303970687, value=1000
              0,00 MAD
 103         column=ventes:produit, timestamp=1613303970510, value=Bure
              aux
 104         column=client:nom, timestamp=1613303970774, value=Othman D
 104         column=client:ville, timestamp=1613303970863, value=Tanger
 104         column=ventes:montant, timestamp=1613303976845, value=1500
              0,00 MAD
 104         column=ventes:produit, timestamp=1613303971052, value=Lits
```

Visualisation de résultat avec la
commande scan

Getting a value :

Get 'name_table', 'num_cell', {COLUMN=>'name_columnfamily:name_column'}

```
hbase(main):020:0> get 'registre_ventes','102',{COLUMN => 'ventes:produit'}
COLUMN          CELL
ventes:produit   timestamp=1613303970001, value=Lampes
1 row(s) in 0.1920 seconds

hbase(main):021:0> disable 'registre_ventes'
0 row(s) in 2.5210 seconds

hbase(main):022:0> drop 'registre_ventes'
0 row(s) in 2.4480 seconds

hbase(main):023:0> exists 'registre_ventes'
Table registre_ventes does not exist
0 row(s) in 0.0550 seconds
```

drop pour supprimer

Exists pour voir vérifier
l'existence de la table

```
agh@agh-VirtualBox:/usr/local/hbase/bin$ cd
agh@agh-VirtualBox:~$ vim .bashrc
agh@agh-VirtualBox:~$ source .bashrc
agh@agh-VirtualBox:~$ cd /home/agh/Downloads
agh@agh-VirtualBox:~/Downloads$ mkdir hbase-code
agh@agh-VirtualBox:~/Downloads$ cd hbase-code/
agh@agh-VirtualBox:~/Downloads/hbase-code$ vim HelloHbase.java
agh@agh-VirtualBox:~/Downloads/hbase-code$
```

```
export CLASSPATH=/usr/local/hbase/lib/*:/usr/local/hadoop/share/hadoop/common/*
-- INSERT --
```

API

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.HColumnDescriptor;
import org.apache.hadoop.hbase.HTableDescriptor;
import org.apache.hadoop.hbase.TableName;
import org.apache.hadoop.hbase.client.*;
import org.apache.hadoop.hbase.util.Bytes;
import java.io.IOException;
public class HelloHBase {
    private Table table1;
    private String tableName = "user";
    private String family1 = "PersonalData";
    private String family2 = "ProfessionalData";
    public void createHbaseTable() throws IOException {
        Configuration config = HBaseConfiguration.create();
        Connection connection = ConnectionFactory.createConnection(config);
        Admin admin = connection.getAdmin();
        HTableDescriptor ht = new HTableDescriptor(TableName.valueOf(tableName));
        ht.addFamily(new HColumnDescriptor(family1));
        ht.addFamily(new HColumnDescriptor(family2));
        System.out.println("connecting");
        System.out.println("Creating Table");
        createOrOverwrite(admin, ht);
        System.out.println("Done.....");
        table1 = connection.getTable(TableName.valueOf(tableName));
        try {
            System.out.println("Adding user: user1");
            byte[] row1 = Bytes.toBytes("user1");
-- INSERT --
```

```
agh@agh-VirtualBox:~/Downloads/hbase-code$ javac HelloHBase.java
agh@agh-VirtualBox:~/Downloads/hbase-code$ ls
HelloHBase.class  HelloHBase.java
agh@agh-VirtualBox:~/Downloads/hbase-code$
```


La table est bien créé , et la lecriture est parfaitement fait , grace aux api en peux manipuler , lire , ecrire les donnes de hbase par un code java aulieu de hbase shell

```
agh@agh-VirtualBox:~/Downloads/hbase-code$ java -cp ./usr/local/hbase/lib/* HelloHBase
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
connecting
Creating Table
Done.....
Adding user: user1
Adding user: user2
reading data...
mohamed
agh@agh-VirtualBox:~/Downloads/hbase-code$
```

Ecriture de fichier dans HDFS

```
agh@agh-VirtualBox:~/Downloads/hbase-code$ cd
agh@agh-VirtualBox:~$ hadoop fs -mkdir -p /input
21/02/14 13:31:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
agh@agh-VirtualBox:~$ hadoop fs -put /home/agh/Downloads/purchases2.txt /input
21/02/14 13:33:54 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
agh@agh-VirtualBox:~$ hadoop fs -ls /input
21/02/14 13:34:51 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 agh supergroup 243309628 2021-02-14 13:34 /input/purchases2.txt
agh@agh-VirtualBox:~$
```



Création de la base de données products

```
hbase(main):001:0> create 'products', 'cf'  
0 row(s) in 9.6590 seconds  
  
=> Hbase::Table - products  
hbase(main):002:0> exit  
agh@agh-VirtualBox:/usr/local/hbase$
```

L'exécution de la commande Importtssv

```
agh@agh-VirtualBox: /usr/local/hbase
completed successfully
2021-02-15 12:32:50,450 INFO [main] mapreduce.Job: Counters: 21
  File System Counters
    FILE: Number of bytes read=60102777
    FILE: Number of bytes written=61275114
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=377535548
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=7
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Map-Reduce Framework
    Map input records=4138476
    Map output records=4138476
    Input split bytes=216
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=112693
    Total committed heap usage (bytes)=170041344
  ImportTsv
    Bad Lines=0
  File Input Format Counters
    Bytes Read=243313724
  File Output Format Counters
    Bytes Written=0
agh@agh-VirtualBox:/usr/local/hbase$
```


L'utilisation de la commande scan avec Limite pour voir une partie de donnees

```
agh@agh-VirtualBox: /usr/local/hbase
scan 'products', {LIMIT=>1}
hbase(main):002:0> scan 'products', {LIMIT=>1}
ROW COLUMN+CELL
1 column=cf:date, timestamp=1613311095452, value=2012-01-01
1 column=cf:payment, timestamp=1613311095452, value=Amex
1 column=cf:price, timestamp=1613311095452, value=214.05
1 column=cf:product, timestamp=1613311095452, value=Men's Clothing
1 column=cf:time, timestamp=1613311095452, value=09:00
1 column=cf:town, timestamp=1613311095452, value=San Jose
1 row(s) in 0.7340 seconds

hbase(main):003:0> scan 'products', {LIMIT=>100}
ROW COLUMN+CELL
1 column=cf:date, timestamp=1613311095452, value=2012-01-01
1 column=cf:payment, timestamp=1613311095452, value=Amex
1 column=cf:price, timestamp=1613311095452, value=214.05
1 column=cf:product, timestamp=1613311095452, value=Men's Clothing
1 column=cf:time, timestamp=1613311095452, value=09:00
1 column=cf:town, timestamp=1613311095452, value=San Jose
10 column=cf:date, timestamp=1613311095452, value=2012-01-01
10 column=cf:payment, timestamp=1613311095452, value=Visa
10 column=cf:price, timestamp=1613311095452, value=213.88
10 column=cf:product, timestamp=1613311095452, value=Toys
10 column=cf:time, timestamp=1613311095452, value=09:00
128844 column=cf:product, timestamp=1613311095452, value=Music
128844 column=cf:time, timestamp=1613311095452, value=12:03
128844 column=cf:town, timestamp=1613311095452, value=Anchorage
128845 column=cf:date, timestamp=1613311095452, value=2012-01-12
128845 column=cf:payment, timestamp=1613311095452, value=Cash
128845 column=cf:price, timestamp=1613311095452, value=460.91
128845 column=cf:product, timestamp=1613311095452, value=Pet Supplies
128845 column=cf:time, timestamp=1613311095452, value=12:03
128845 column=cf:town, timestamp=1613311095452, value=Jersey City
128846 column=cf:date, timestamp=1613311095452, value=2012-01-12
128846 column=cf:payment, timestamp=1613311095452, value=Discover
128846 column=cf:price, timestamp=1613311095452, value=440.16
128846 column=cf:product, timestamp=1613311095452, value=Pet Supplies
128846 column=cf:time, timestamp=1613311095452, value=12:03
128846 column=cf:town, timestamp=1613311095452, value=Charlotte
128847 column=cf:date, timestamp=1613311095452, value=2012-01-12
128847 column=cf:payment, timestamp=1613311095452, value=MasterCard
128847 column=cf:price, timestamp=1613311095452, value=30.9
128847 column=cf:product, timestamp=1613311095452, value=Baby
128847 column=cf:time, timestamp=1613311095452, value=12:03
128847 column=cf:town, timestamp=1613311095452, value=Virginia Beach
128848 column=cf:date, timestamp=1613311095452, value=2012-01-12
128848 column=cf:payment, timestamp=1613311095452, value=MasterCard
128848 column=cf:price, timestamp=1613311095452, value=257.63
```

Lire une ligne

```
hbase(main):004:0> get 'products','2000',{COLUMN => 'cf:town'}  
COLUMN          CELL  
cf:town          timestamp=1613311095452, value=Oklahoma City  
1 row(s) in 0.1810 seconds  
  
hbase(main):005:0>
```

Traitement de données avec Spark :

```
agh@agh-VirtualBox:/usr/local/hbase$ cd ~/Downloads
agh@agh-VirtualBox:~/Downloads$ ls
apache-maven-3.5.0-bin.tar.gz  jdk-8u71-linux-x64.tar.gz
hadoop-2.7.4.tar.gz           purchases2.txt
hbase-1.4.7-bin.tar.gz       spark-2.4.3-bin-hadoop2.7.tgz
hbase-code                    tp_hbase.pdf
agh@agh-VirtualBox:~/Downloads$ tar -zxvf apache-maven-3.5.0-bin.tar.gz
apache-maven-3.5.0/README.txt
apache-maven-3.5.0/LICENSE
apache-maven-3.5.0/NOTICE
apache-maven-3.5.0/lib/

agh@agh-VirtualBox:~/Downloads$ sudo mv /home/agh/Downloads/apache-maven-3.5.0
/opt/
[sudo] password for agh:
agh@agh-VirtualBox:~/Downloads$ ls
apache-maven-3.5.0-bin.tar.gz  jdk-8u71-linux-x64.tar.gz
hadoop-2.7.4.tar.gz           purchases2.txt
hbase-1.4.7-bin.tar.gz       spark-2.4.3-bin-hadoop2.7.tgz
hbase-code                    tp_hbase.pdf
agh@agh-VirtualBox:~/Downloads$


agh@agh-VirtualBox:~/Downloads$ sudo vim /etc/profile
agh@agh-VirtualBox:~/Downloads$ source /etc/profile
agh@agh-VirtualBox:~/Downloads$ mvn -v
Apache Maven 3.5.0 (ff8f5e7444045639af65f6095c62210b5713f426; 2017-04-03T20:39:
06+01:00)
Maven home: /opt/apache-maven-3.5.0
Java version: 1.8.0_71, vendor: Oracle Corporation
Java home: /opt/java/jdk1.8.0_71/jre
Default locale: en_US, platform encoding: UTF-8
OS name: "linux", version: "5.8.0-43-generic", arch: "amd64", family: "unix"
agh@agh-VirtualBox:~/Downloads$
```

Installation de maven

Création de projet maven Myapp

```
----  
[INFO] Parameter: basedir, Value: /home/agh  
[INFO] Parameter: package, Value: TP.Myapp  
[INFO] Parameter: groupId, Value: TP.Myapp  
[INFO] Parameter: artifactId, Value: MyApp  
[INFO] Parameter: packageName, Value: TP.Myapp  
[INFO] Parameter: version, Value: 1.0-SNAPSHOT  
[INFO] project created from Old (1.x) Archetype in dir: /home/agh/Myapp  
[INFO] -----  
[INFO] BUILD SUCCESS  
[INFO] -----  
[INFO] Total time: 01:01 min  
[INFO] Finished at: 2021-02-15T12:43:00Z  
[INFO] Final Memory: 15M/44M  
[INFO] -----  
agh@agh-VirtualBox:~$
```

```
agh@agh-VirtualBox:~/Myapp$ mvn package  
[INFO] Scanning for projects...
```



```
[INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ Myapp ---  
[INFO] Building jar: /home/agh/Myapp/target/Myapp-1.0-SNAPSHOT.jar  
[INFO] -----  
[INFO] BUILD SUCCESS  
[INFO] -----  
[INFO] Total time: 11.876 s  
[INFO] Finished at: 2021-02-15T12:44:56Z  
[INFO] Final Memory: 15M/44M  
[INFO] -----  
agh@agh-VirtualBox:~/Myapp$
```

Configuration de projet maven : ajouter les dépendences (spark,hbase) de notre projet

```
agh@agh-VirtualBox: ~/Myapp
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>

  <groupId>TP.Myapp</groupId>
  <artifactId>Myapp</artifactId>
  <version>1.0-SNAPSHOT</version>
  <packaging>jar</packaging>

  <name>Myapp</name>
  <url>http://maven.apache.org</url>

  <properties>
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
  </properties>
  <build>
</build>
<plugins>
<plugin>
<groupId>org.apache.maven.plugins</groupId>
<artifactId>maven-compiler-plugin</artifactId>
<configuration>
<source>1.8</source>
<target>1.8</target>
</configuration>
</plugin>
-- INSERT --
```

```
agh@agh-VirtualBox: ~/Myapp
</plugin>
</plugins>
</build>
  <dependencies>
    <dependency>
      <groupId>junit</groupId>
      <artifactId>junit</artifactId>
      <version>3.8.1</version>
      <scope>test</scope>
    </dependency>
    <dependency>
      <groupId>org.apache.hbase</groupId>
      <artifactId>hbase</artifactId>
      <version>2.1.3</version>
      <type>pom</type>
    </dependency>
    <dependency>
      <groupId>org.apache.hbase</groupId>
      <artifactId>hbase-spark</artifactId>
      <version>2.0.0-alpha4</version>
    </dependency>
    <dependency>
      <groupId>org.apache.spark</groupId>
      <artifactId>spark-core_2.11</artifactId>
      <version>2.2.1</version>
    </dependency>
  </dependencies>
</project>
-- INSERT --
```

```
agh@agh-VirtualBox:~/Myapp/src/main/java/TP/Myapp$ ls
App.java
agh@agh-VirtualBox:~/Myapp/src/main/java/TP/Myapp$ mv App.java
mv: missing destination file operand after 'App.java'
Try 'mv --help' for more information.
agh@agh-VirtualBox:~/Myapp/src/main/java/TP/Myapp$ mv App.java HbaseSparkProcess.java
agh@agh-VirtualBox:~/Myapp/src/main/java/TP/Myapp$ ls
HbaseSparkProcess.java
agh@agh-VirtualBox:~/Myapp/src/main/java/TP/Myapp$
```

```
[INFO]
[INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ Myapp ---
[INFO] Building jar: /home/agh/Myapp/target/Myapp-1.0-SNAPSHOT.jar
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 12.865 s
[INFO] Finished at: 2021-02-15T13:00:15Z
[INFO] Final Memory: 44M/105M
[INFO] -----
```

```
agh@agh-VirtualBox:~/Myapp$
```

```
[INFO]
[INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ Myapp ---
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 13.401 s
[INFO] Finished at: 2021-02-15T13:01:34Z
[INFO] Final Memory: 44M/106M
[INFO] -----
```

```
agh@agh-VirtualBox:~/Myapp$
```


Le résultat de compter le nombre d'élément



```
nombre d'enregistrements: 4138476 ✓
21/02/15 13:07:28 INFO SparkContext: Invoking stop() from shutdown hook
21/02/15 13:07:29 INFO SparkUI: Stopped Spark web UI at http://10.0.2.15:4040
21/02/15 13:07:29 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
21/02/15 13:07:29 INFO MemoryStore: MemoryStore cleared
21/02/15 13:07:29 INFO BlockManager: BlockManager stopped
21/02/15 13:07:29 INFO BlockManagerMaster: BlockManagerMaster stopped
21/02/15 13:07:29 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
21/02/15 13:07:29 INFO SparkContext: Successfully stopped SparkContext
21/02/15 13:07:29 INFO ShutdownHookManager: Shutdown hook called
21/02/15 13:07:29 INFO ShutdownHookManager: Deleting directory /tmp/spark-46c3dbb2-8739-43cd-8772-3c015169e138
21/02/15 13:07:29 INFO ShutdownHookManager: Deleting directory /tmp/spark-9d16fe00-e410-45ae-917f-108e64d6225a
agh@agh-VirtualBox:~/Myapp/target$
```

Activer V



• **Merci**